

# $\mu$ TP

Heze Yin, Laixian Wan, Peiqing Lu, Xiaoduan Chang

GitHub link: [https://github.com/drcxd/CS655\\_GENI\\_Project\\_uTP](https://github.com/drcxd/CS655_GENI_Project_uTP)

Project on GENI: `utp-experiment`

## 1 Introduction

One major problem that BitTorrent users experience is throttling of non-BitTorrent TCP connections. A BitTorrent client can simultaneously talk to many peers, each using multiple TCP connections. As a result, non-BitTorrent TCP connections are limited in speed, sometimes significantly when there are tens of BitTorrent TCP connections. This is because TCP protocol is designed to split bandwidth equally among all TCP connections. To overcome this problem, Micro Transport Protocol, also known as  $\mu$ TP or uTP is developed. It is an open-sourced UDP-based variant of the BitTorrent peer-to-peer file sharing protocol intends to mitigate poor latency and other congestion control problems found in conventional BitTorrent over TCP but still provide reliable, ordered delivery.

In this project, we are going to investigate into  $\mu$ TP's mechanism and also conduct experiments to evaluate how effective  $\mu$ TP can be, comparing to traditional BitTorrent file sharing protocol over TCP.

Through this experiment, we have learned about mechanism of  $\mu$ TP, more detail about how BitTorrent protocol works and also several tools which could help ease the process of experiment.

## 2 Experimental Methodology

First of all, we need to verify the problem of traditional BitTorrent file sharing problem over TCP, i.e. throttling of other non-BitTorrent TCP connections. For this purpose, we need to find a BitTorrent client which has not implement  $\mu$ TP. After some investigation, we decide to use **aria2**<sup>1</sup>, which is well-documented and provides a easy to use command-line interface. This makes it a good choice to use on GENI virtual machines.

Then, to validate the effectiveness of  $\mu$ TP and compare with BitTorrent over TCP, we have to use another BitTorrent client which has implemented  $\mu$ TP. **qBittorrent** was our initial choice, however, it turns out to be difficult to use only under command-line interface. Thus, we choose **transmission**<sup>2</sup>. To be specific, its command-line variant **transmission-daemon** along with **transmission-remote** to control the daemon process.

### 2.1 Slice Topology

The topology of the GENI slice we use is shown in the following figure:




Figure 1: Topology

In this slice we have four nodes or hosts (we will use the term "node" and "host" interchangeably to refer to virtual machines in the network) connected by one link.

**node-0** and **node-1** are the seeding nodes. They possess the whole file in their local hard drive and transmit data to downloading nodes. **node-2** is the single downloading node. It will request file from the two seeding nodes using BitTorrent protocol. **node-3** is the testing node. It use **iperf** to establish a TCP connection with **node-2** and measure the throughput.

## 2.2 Limit Bandwidth

To get reasonable experiment result, we have to put some constraints on the bandwidth of network interfaces of these nodes.

First, to have enough time to operate, we have to either transmit a extremely large file, e.g. 20GB, or limit the speed of the BitTorrent client. We choose the later since it looks more flexible. To be specific, we limit the upload bandwidth of two seeding nodes to 1024KB/s.

Second, the BitTorrent client download bit rate should greater than  $2/3$  of the download bandwidth, otherwise, we may not observe the throttling of other TCP connections. In our case, there is only one other TCP connection established by `iperf` server and client. Assume upload bit rate is  $u$ , then the download bit rate  $d$  of the downloading node should be in the range of  $(1u, 3u)$ . Thus, we limit the download bandwidth of downloading node to 2048KB/s.

No constraints need to be put on the testing node.

## 2.3 Expected Results

Assume the download bandwidth is  $d$ , then when using BitTorrent over TCP, we expect the throughput measured by `iperf` be around  $d/3$ ; when using  $\mu$ TP, we expect the throughput measured by `iperf` be at least  $d/2$ .

# 3 Results

## 3.1 Usage Instructions

Currently the four nodes are ready to conduct the experiment. We can login and follow the instructions below.

Since the configuration for the BitTorrent clients are local to users, we need to switch to user `xiaoduan` to execute the current commands correctly:

```
sudo su xiaoduan
```

### 3.1.1 Common Setup

#### 1. Launch the Tracker

Each BitTorrent client could communicate with each other with the help of a tracker service. The tracker is deployed on `node-0`. Before doing any of the two experiments, we have to launch the tracer first. On `node-0`, execute:

```
# use another screen to run the tracker, it's like open another
# instance of shell
screen
bttrack --port 6969 --dfile ~/.bttrack/dstate --logfile ~/.bttrack/tracker.log \
--nat_check 0 --scrape_allowed full
# The tracker launches and output nothing. Press Ctrl-A + d to switch
# back to the original shell
```

#### 2. Limit Bandwidth

We also have to limit the bandwidth of `nodes-0`, `nodes-1` and `nodes-2` before any experiment.

On `node-0` and `node-1` execute:

```
# This limit the upload bandwidth to 1024KB/s
cd ~
sudo wondershaper/wondershaper -a eth1 -u 1024
# To clear all constraints on an interface
# sudo wondershaper/wondershaper -ca eth1
```

On `node-2` execute:

```
# This limit the download bandwidth to 2048KB/s
cd ~
sudo wondershaper/wondershaper -a eth1 -d 2048
```

#### 3. Launch iperf Server on node-2

On `node-2`:

```
screen
iperf -s
# Use Ctrl-A + d to return to the original shell
```

### 3.1.2 Using aria2 without $\mu$ TP

Seeding on node-0 and node-1:

```
cd ~
aria2c -V beethoven.torrent
```

Start downloading on node-2:

```
# make sure that the downloads directory is empty
cd ~
rm downloads/*
aria2c beethoven.torrent
```

Start iperf client on node-3:

```
# Set a longer time, so all three connections could reach stable state
# We may have to terminate it using Ctrl-C manually when the
# downloading finishes. Otherwise, the average throughput will count
# throughput when there is no BitTorrent connections.
iperf -c 10.10.1.3 -t 180
```

We could stop one or both of the seeding processes to test the throughput when there is only 1 or 0 BitTorrent TCP connection.

### 3.1.3 Using transmission with $\mu$ TP

Using transmission is more complicated, because we have to launch a daemon process and use another program to control it.

Make sure the tracker and iperf server are running and constraints are still on interfaces.

Seeding on node-0 and node-1:

```
screen # run daemon in a new shell
transmission-daemon -f # force transmission-daemon in the foreground
# Use Ctrl-A + d to return to the original shell
cd ~;
transmission-remote -a beethoven.torrent
# show all torrents
# transmission-remote -l
# remove torrent with torrent-id, torrent-id is displayed in the
# output of transmission-remote -l
# transmission-remote -t torrent_id -r
```

One thing worth noting is that transmission remembers all the torrents once they are added, so if we have added a torrent once, we don't have to add it again. Also, if we want to download a torrent once again, we have to remove both the torrent from transmission and the local file (not the .torrent file).

Start downloading on node-2:

```
screen
transmission-daemon -f
# Ctrl-A + d return to original shell
cd ~;
rm downloads/*
# remove all torrents if necessary
# transmission-remote -t all -r
transmission-remote -a beethoven.torrent
# display downloading information, including speed
# transmission-remote -l
```

Start iperf client on node-3:

```
iperf -c 10.10.1.3 -t 180
```

We could test throughput when there is only 1 or 0 peer by closing one or both transmission-daemon processes on node-0 and node-1.

### 3.1.4 Cleanup

Remember to stop all the aria2c processes, transmission-daemon processes, screen sessions, iperf server process and the tracker process. Also, clear all the constraints we put on the network interfaces.

## 3.2 Analysis

### 3.2.1 Using aria2 without $\mu$ TP

The output of iperf client detecting the throughput when there are two other BitTorrent TCP connections:

```
xiaoduan@node-3:~$ iperf -c 10.10.1.3 -t 180
-----
Client connecting to 10.10.1.3, TCP port 5001
TCP window size: 85.0 KByte (default)
-----
[ 3] local 10.10.1.4 port 37852 connected with 10.10.1.3 port 5001
^C[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-165.0 sec 13.6 MBytes 693 Kbits/sec
```

The output of iperf client detecting the throughput when there is one BitTorrent TCP connection:

```
xiaoduan@node-3:~$ iperf -c 10.10.1.3 -t 180
-----
Client connecting to 10.10.1.3, TCP port 5001
TCP window size: 85.0 KByte (default)
-----
[ 3] local 10.10.1.4 port 37854 connected with 10.10.1.3 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-181.1 sec 26.6 MBytes 1.23 Mbits/sec
```

The output of iperf client detecting the throughput when there is no other TCP connection:

```
xiaoduan@node-3:~$ iperf -c 10.10.1.3 -t 30
-----
Client connecting to 10.10.1.3, TCP port 5001
TCP window size: 85.0 KByte (default)
-----
[ 3] local 10.10.1.4 port 37856 connected with 10.10.1.3 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-30.9 sec 7.62 MBytes 2.07 Mbits/sec
```

The experiment result meets our expectation that when there are two BitTorrent TCP connections, the iperf TCP connection only have around 1/3 share of the total bandwidth.

### 3.2.2 Using transmission with $\mu$ TP

The output of iperf client detecting the throughput when there is two peers connected:

```
xiaoduan@node-3:~$ iperf -c 10.10.1.3 -t 180
-----
Client connecting to 10.10.1.3, TCP port 5001
TCP window size: 85.0 KByte (default)
-----
[ 3] local 10.10.1.4 port 37862 connected with 10.10.1.3 port 5001
^C[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-83.0 sec 17.9 MBytes 1.81 Mbits/sec
```

When iperf is running, the output of transmission-remote -l on node-2:

```
xiaoduan@node-2:~$ transmission-remote -l
ID      Done      Have  ETA      Up    Down  Ratio  Status      Name
  2     94%    25.09 MB 12 sec    0.0   26.0    0.0  Downloading Beethoven_Symphony_N05_IV_Finale.mp3
Sum:           25.09 MB          0.0   26.0
```

From the result we can see that the  $\mu$ TP implementation of transmission reduces its own bit rate significantly when there is other TCP connections.

Here is also a figure 2 shows how the download speed changes. We launch iperf client for 30 seconds multiple time during the whole process. Every time the download speed drops drastically, we know that the iperf client is launched.

## 4 Conclusion

From our experiments, we confirm that BitTorrent protocol implemented over TCP does throttle other non-BitTorrent connections, especially when there are multiple BitTorrent TCP connections.

On the other hand,  $\mu$ TP implemented by transimisson over UDP is much less aggressive. In fact, it almost stop downloading entirely when there is other competing TCP connection. Anyway, we can see that  $\mu$ TP implemented over UDP

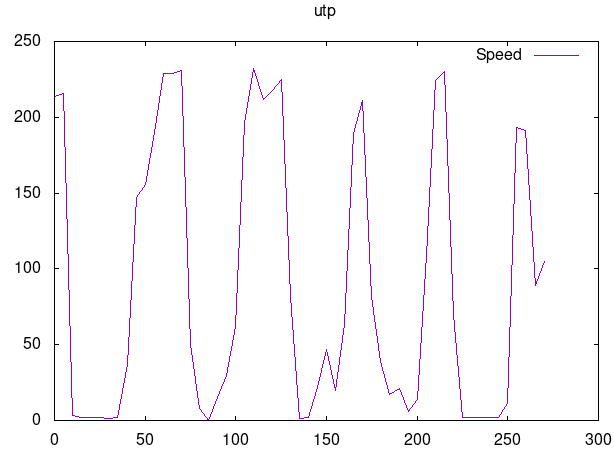


Figure 2:  $\mu$ TP

does address the issue while still provides reliable and ordered delivery. It does achieve the goal that BitTorrent clients using  $\mu$ TP do not disturb other internet connections while still utilizing the unused bandwidth fully. In our case, since `iperf` is a bandwidth measurement program and tries to take as much bandwidth as possible, the download speed of BitTorrent clients drops to almost 0.

To extend our experiment of  $\mu$ TP, we could also test its behavior when it works with other TCP connections which are not as aggressive as the `iperf` and see how it adjusts BitTorrent client's download speed instead of almost stopping at all.

## 5 Division of Labor

Xiaoduan Cheng: configured the environment, designed the experiment procedure, wrote the report.

Heze Yin: did the experiment, recorded the demo video, wrote the report.

Laixian Wan: did the experiment, wrote the report.

Peiqing Lu: did the experiment, wrote the report.

## 6 Reference

1. aria2 <https://aria2.github.io/>
2. transmission <https://transmissionbt.com/>
3. Wikipedia [https://en.wikipedia.org/wiki/Micro\\_Transport\\_Protocol](https://en.wikipedia.org/wiki/Micro_Transport_Protocol)
4. BEP-29 [http://www.bittorrent.org/beps/bep\\_0029.html](http://www.bittorrent.org/beps/bep_0029.html)
5. RFC 6817 <https://tools.ietf.org/html/rfc6817>